

Operating Systems

Lecture 01

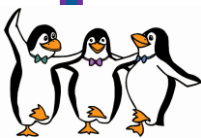
Introduction to OS

Dr. Khalid A. Hafeez



Objectives

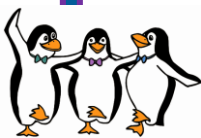
- To describe the basic organization of computer systems
- To provide a grand tour of the major components of an operating systems
- To give an overview of the many types of computing environments
- To explore several open-source operating systems





Introduction

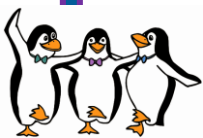

- What is an Operating System?
 - A program that acts as an intermediary between a user of a computer and the computer hardware
 - Operating system goals:
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner
- *Some operating systems are designed to be convenient, others to be efficient, and others to be some combination of the two.*





Computer System Structure


- Computer system can be divided into four components:





Computer System Structure


- Computer system can be divided into four components:
 - **Hardware** – provides basic computing resources
 - CPU, memory, I/O devices
 - **Operating system**
 - Controls and coordinates use of hardware among various applications and users
 - **Application programs** – define the ways in which the system resources are used to solve the computing problems of the users
 - Word processors, compilers, web browsers, database systems, video games
 - **Users**
 - People, machines, other computers





What Operating Systems Do

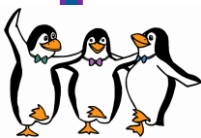
- Exploring the OSES from two viewpoints:
 - **User view:**
 - Users want convenience, **ease of use** and **good performance**
 - Don't care about **resource utilization**
 - But shared computer such as **mainframe** or **minicomputer** must keep all users happy
 - Users of dedicate systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
 - Handheld computers (smartphones) are resource poor, optimized for usability and battery life
 - Some computers have little or no user interface, such as embedded computers in devices and automobiles





What Operating Systems Do


- Exploring the OSES from two viewpoints:
 - **System view:**
 - The OS is the program that most intimately involved with the hardware
 - We can view the OS as a:
 - **Resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
 - **Control program**
 - Controls the execution of user programs to prevent errors and improper use of the computer





What is Operating System?


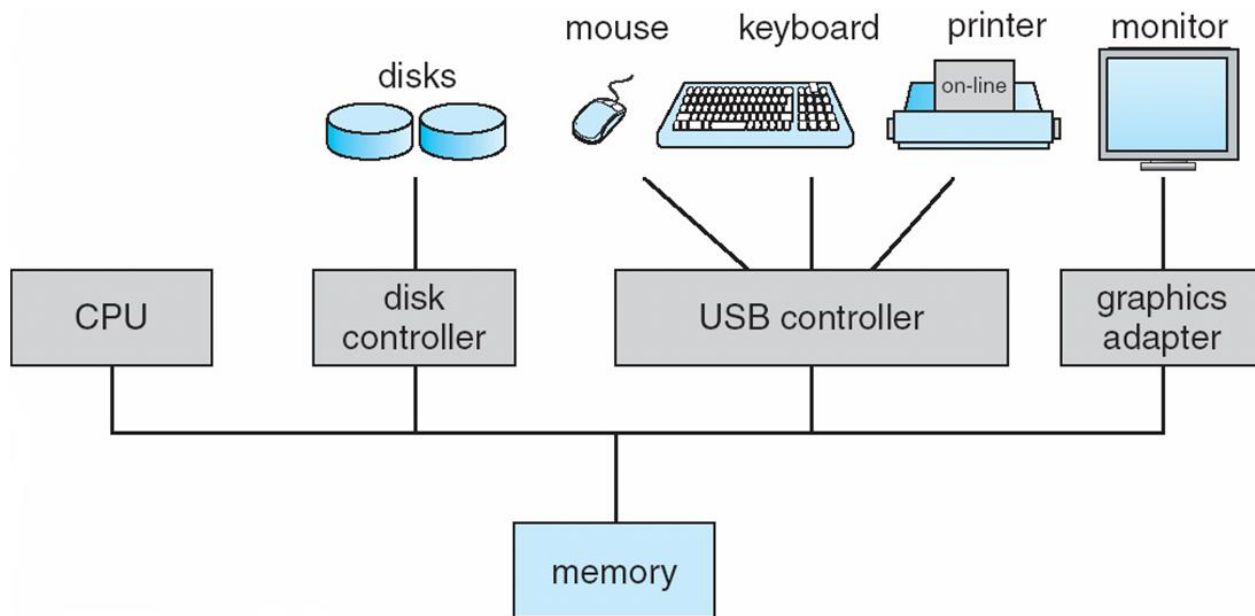
- Operating System Definition:
 - No universally accepted definition
 - “Everything a vendor ships when you order an operating system” is a good approximation, but varies wildly
 - Computer hardware is constructed toward executing user programs.
 - Since bare hardware alone is not easy to use, application programs are developed.
 - These programs require certain common operations, such as those controlling the I/O devices.
 - The common functions of controlling and allocating resources are then brought together into one piece of software: the **operating system**.
 - Common definition:
 - “The one program running at all times on the computer” is the **kernel**.
 - And two other types of programs:
 - a system program (ships with the operating system) ,
 - an application program.





Computer System Organization

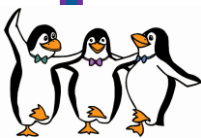
- Computer System Operation:
 - One or more CPUs, device controllers connected through common bus that provide access to shared memory
 - The CPU and the device controllers can execute in parallel, competing for memory cycles





Computer System Organization

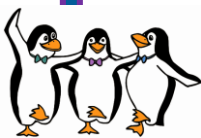
- Computer System Operation:
 - I/O devices and the CPU can execute concurrently
 - Each device controller is in charge of a particular device type
 - Each device controller has a local buffer
 - CPU moves data from/to main memory to/from local buffers
 - I/O is from the device to local buffer of controller
 - Device controller informs CPU that it has finished its operation by causing an **interrupt**





Computer System Organization

- Computer System Operation:
 - For a computer to start running, it needs an initial program to run
 - **bootstrap program** is loaded at power-up or reboot
 - Typically stored in **ROM** or **EEPROM**, generally known as **firmware**
 - Initializes all aspects of system
 - Loads operating system **kernel** and starts execution
 - The kernel then starts providing services to the system and its users.
 - Some services are provided outside of the kernel, by system programs that are loaded into memory at boot time to become **system processes**, or **system daemons** that run the entire time the kernel is running.
 - The system is now running and waiting for an event to occur






Computer System Organization

- Common Functions of Interrupts:


- The occurrence of an event is usually signaled by an **interrupt** from either the hardware or the software:
 - Hardware trigger an interrupt by sending a **signal** to the CPU
 - Software trigger an interrupt by executing a special operation called a **system call**
- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- An operating system is **interrupt driven**






Computer System Organization

- Common Functions of Interrupts:
 - Interrupt timeline




Interrupt timeline for a single process doing output.





Computer System Organization

- Interrupt Handling:
 - The operating system preserves the state of the CPU by storing registers and the program counter
 - Determines which type of interrupt has occurred:
 - **polling**
 - **vectored** interrupt system
 - Separate segments of code determine what action should be taken for each type of interrupt

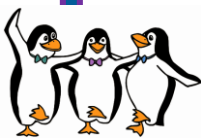




Computer System Organization

- Storage Structure:

- **Main memory:** only large storage media that the CPU can access directly
 - **Random access**
 - Typically **volatile**
- **Secondary storage:** extension of main memory that provides large **nonvolatile** storage capacity
- **Hard disks:** rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer
- **Solid-state disks:** faster than hard disks, nonvolatile
 - Various technologies
 - Becoming more popular


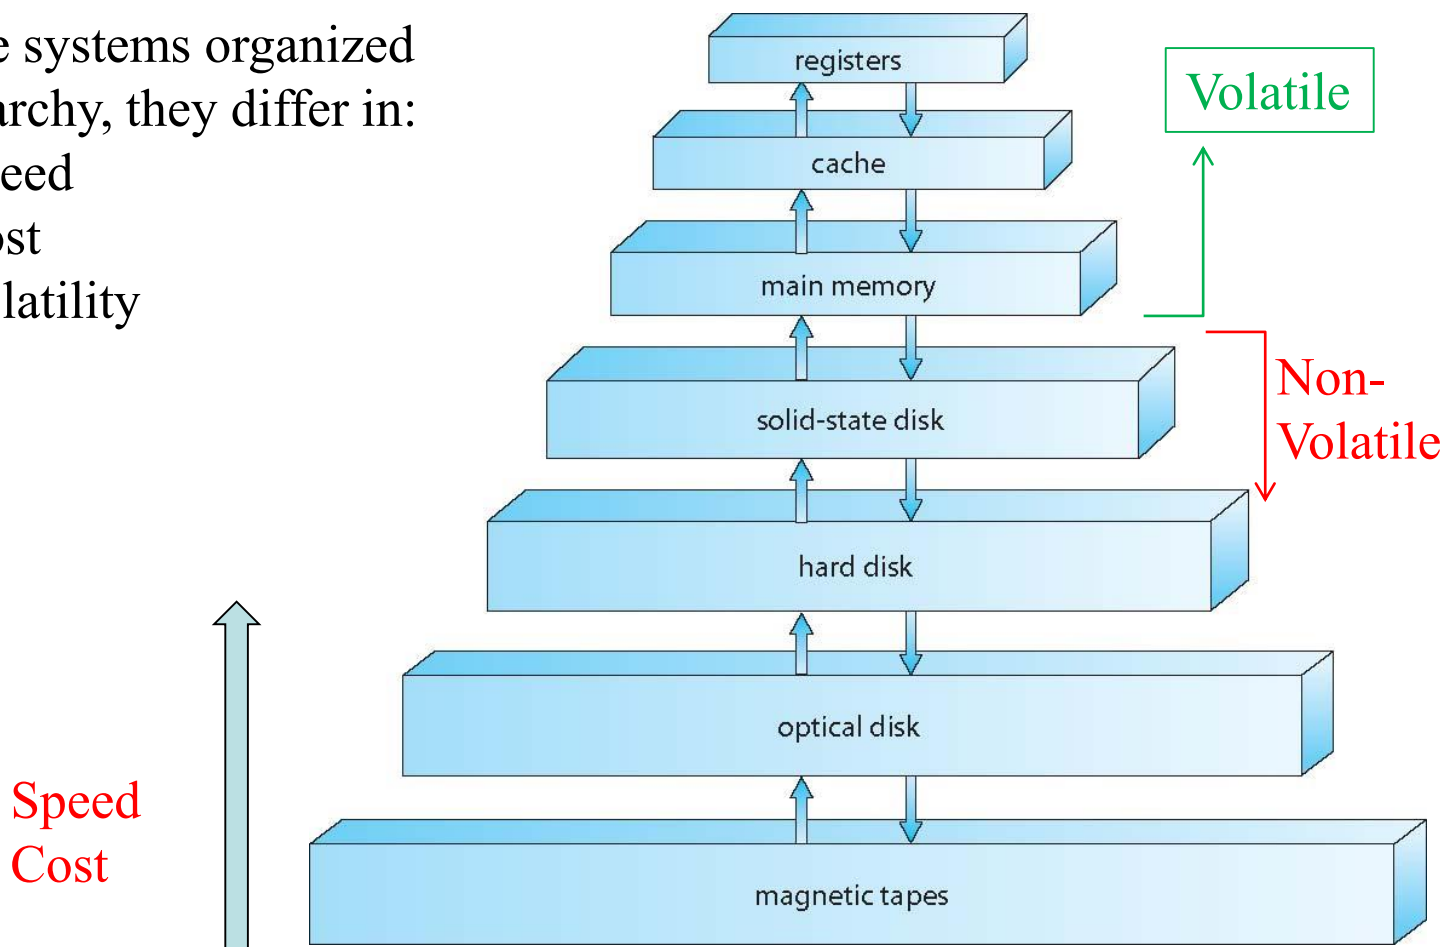




Computer System Organization

- Storage Hierarchy:

- Storage systems can be organized in a hierarchy according to speed and cost.
- Storage systems organized in hierarchy, they differ in:
 - Speed
 - Cost
 - Volatility

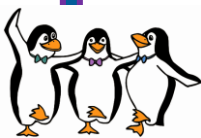




Computer System Organization

- Caching:


- **Caching**: copying information into faster storage system; main memory can be viewed as a cache for secondary storage
- It is an important principle, that is performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management is an important design problem
 - Cache size and replacement policy





Computer System Organization

- Storage Definitions and Notation Review:
 - The basic unit of computer storage is the **bit**.
 - A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage.
 - A **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes.
 - Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes.
 - **kilobyte**, or **KB**, is 1,024 bytes (2^{10})
 - **megabyte**, or **MB**, is 1,024² bytes (2^{20})
 - **gigabyte**, or **GB**, is 1,024³ bytes (2^{30})
 - **terabyte**, or **TB**, is 1,024⁴ bytes (2^{40})
 - **petabyte**, or **PB**, is 1,024⁵ bytes (2^{50})

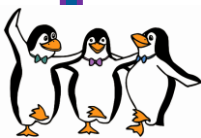




Computer System Organization

- I/O Structure:


- A general-purpose computer system consists of CPUs and multiple device controllers that are connected through a common bus.
- Each device controller is in charge of a specific type of device.
- Operating systems have a **device driver** for each device controller to manage I/O





Computer System Organization

- I/O Structure:
 - **Interrupt-driven I/O** is good for moving small amounts of data
 - The device driver loads the appropriate registers within the device controller.
 - The device controller determines what action to take (such as “read a character from the keyboard”).
 - The controller transfers the data from the device to its local buffer.
 - Then the device controller informs the device driver via an interrupt that it has finished its operation.
 - The device driver then returns control to the operating system
 - This can produce high overhead when used for bulk data movement
 - To solve this problem, **direct memory access (DMA)** is used


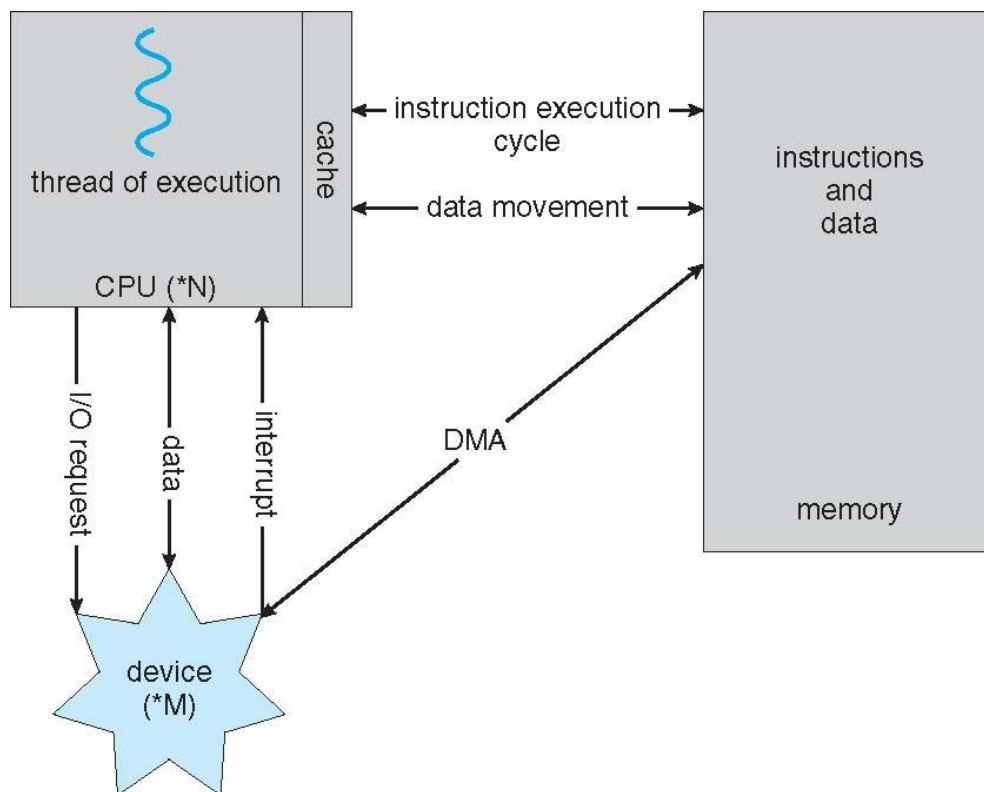




Computer System Organization

- Direct Memory Access Structure:


- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte





Computer-System Architecture

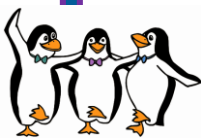

- **Single-Processor Systems:**
 - Most systems use a single general-purpose processor
 - Most systems have special-purpose processors as well
 - For keyboard, disks, ...
- **Multiprocessors** systems growing in use and importance
 - Also known as **parallel systems**, **tightly-coupled systems**
 - They have two or more processors in close communication, sharing the computer resources
 - Advantages include:
 1. **Increased throughput:** more work done in less time
 2. **Economy of scale:** cost less than equivalent multiple single-processor systems
 3. **Increased reliability:** graceful degradation or fault tolerance





Computer-System Architecture

- The multiple-processor systems in use today are of two types:
 - **Asymmetric Multiprocessing**: each processor is assigned a specific task.
 - This scheme defines a **boss-worker** relationship.
 - The boss processor schedules and allocates work to the worker processors.
 - **Symmetric Multiprocessing (SMP)**:
 - It is the most common
 - All processors are peers: each processor performs all tasks
 - Symmetric Multiprocessing Architecture:







Computer-System Architecture

- Multi-Core CPUs:

- They are more efficient than multiple chips with single cores
 - Because on-chip communication is faster than between-chip communication.
- Uses significantly less power than multiple single-core chips
- These multicore CPUs appear to the operating system as N standard processors.
- dual-core design with two cores on the same chip.

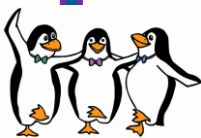





Computer-System Architecture

- Clustered Systems:

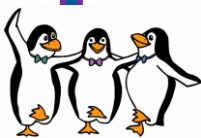

- They are like multiprocessor systems, but multiple systems working together
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service which survives failures
 - **Asymmetric clustering** has one machine in hot-standby mode
 - **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - Some clusters are for **high-performance computing (HPC)**
 - Applications must be written to use **parallelization**
 - Some have **distributed lock manager (DLM)** to avoid conflicting operations





Operating System Structure

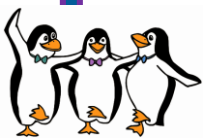
- Important Aspects of Operating Systems:
 - **Multiprogramming (Batch system)** is needed for efficiency:
 - Single user cannot keep CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - A subset of total jobs in system is kept in memory
 - One job selected and run via **job scheduling**
 - When it has to wait (for I/O for example), OS switches to another job





Operating System Structure


- Important Aspects of Operating Systems:
 - **Timesharing (multitasking)** is logical extension to multiprogramming in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - **Response time** should be < 1 second
 - Each user has at least one program executing in memory \Rightarrow **process**
 - If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out to run
 - **Virtual memory** allows execution of processes that are larger than actual physical memory





Operating-System Operations

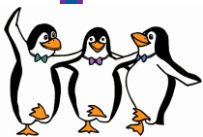
- Modern operating systems are interrupt driven:
 - If there are no processes, no I/O devices, and no users to whom to respond, an operating system will sit quietly.
 - **Interrupt driven** by hardware or software
 - Hardware interrupt by one of the devices
 - Software interrupt (called **exception** or **trap**):
 - Software error (e.g., division by zero)
 - Request for operating system service
 - Other process problems include infinite loop, processes modifying each other or even the operating system itself





Operating-System Operations


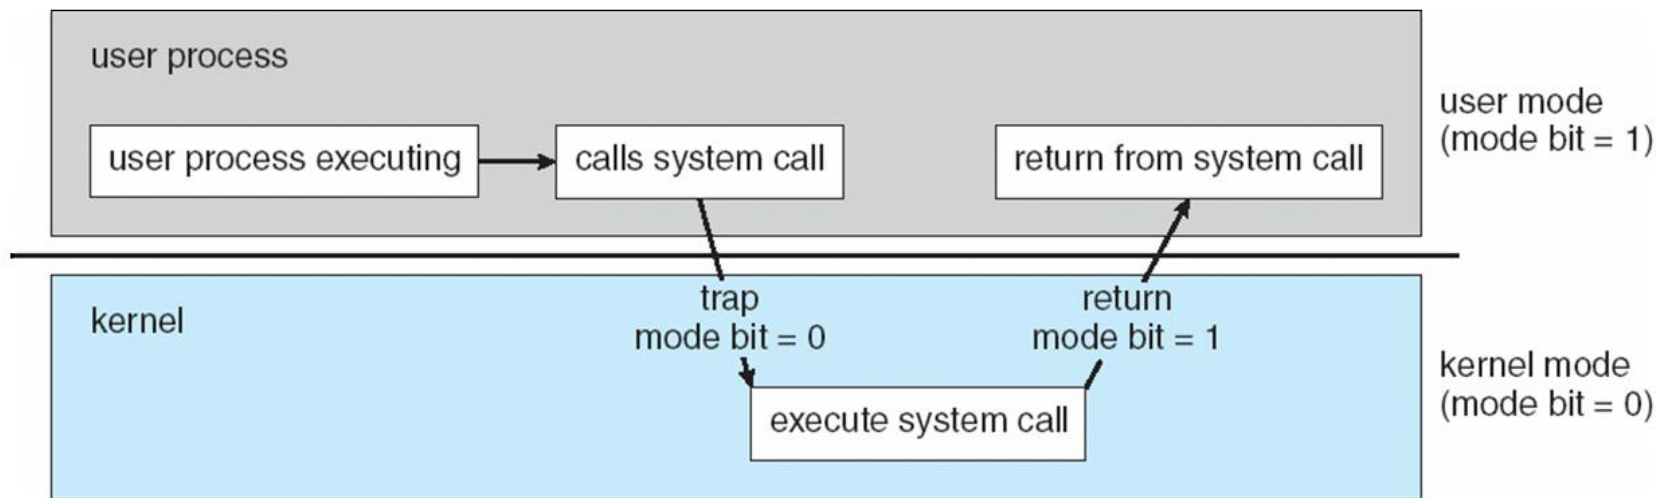
- Dual-Mode and Multi-Mode Operation:
 - To ensure the proper execution of the OS, we must be able to distinguish between the execution of operating-system code and user-defined code.
 - Therefore, computer systems provide hardware bit to differentiate among various modes of execution.
 - **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode** (called **supervisor**, **system**, or **privileged mode**)
 - **Mode bit** provided by hardware (kernel (0), user (1))
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode (such as: switch mode, I/O control, timer management, and interrupt management instructions)





Operating-System Operations

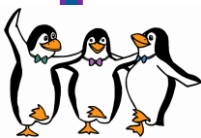
- Dual-Mode and Multi-Mode Operation:
 - Transition from user to kernel mode:
 - **System call** changes mode to kernel, return from call resets it to user





Operating-System Operations


- Dual-Mode and Multi-Mode Operation:
 - The concept of modes can be extended beyond two modes
 - Increasingly CPUs support multi-mode operations
 - CPU uses more than one bit to set and test the mode.
 - Example:
 - CPUs that support virtualization frequently, have a separate mode to indicate when the **virtual machine manager (VMM)** and the virtualization management software is in control of the system.
 - In this mode, the VMM has more privileges than user processes but fewer than the kernel





Operating-System Operations


- Transition from User to Kernel Mode:
 - To ensure that the OS maintains control over the CPU, we use timers.
 - **Timer** is used to prevent a user program to get stuck in infinite loop / process hogging resources
 - Timer is set to interrupt the computer after some time period
 - Keep a counter that is decremented by the physical clock.
 - Operating system set the counter (it is a privileged instruction)
 - When counter reaches zero, it generates an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time





Process Management

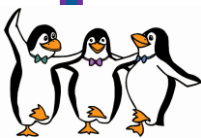
- Process Management:
 - A **process** is a program in execution. It is a unit of work within the system. Program is a *passive entity*, process is an *active entity*.
 - Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
 - Process termination requires reclaim of any reusable resources
 - **Single-threaded** process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
 - **Multi-threaded** process has one program counter per thread
 - Typically, a system has many processes, some user, some operating system running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes / threads





Process Management

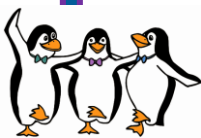
- Process Management Activities:
 - The operating system is responsible for the following activities in connection with process management:
 - Scheduling processes and threads on the CPUs
 - Creating and deleting both user and system processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling





Memory Management


- Memory Management:
 - To execute a program all (or part) of the instructions must be in memory
 - All (or part) of the data that is needed by the program must be in memory.
 - Memory management determines what is in memory and when
 - Optimizing CPU utilization and computer response to users
 - Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed





Storage Management

- Storage Management:
 - OS provides uniform, logical view of information storage
 - It abstracts physical properties to logical storage unit - **file**
 - Each medium is controlled by device (i.e., disk drive, tape drive)
 - Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
 - File-System management
 - It is one of the most visible components of an operating system.
 - Files usually organized into directories
 - Access control on most systems to determine who can access what
 - OS activities include:
 - Creating and deleting files and directories
 - Support primitives to manipulate files and directories
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media






Storage Management

- Mass-Storage Management:

- Usually disks used to store data that does not fit in main memory or data that must be kept for a “long” period of time
- Proper management is of central importance
- Entire speed of computer operation hinges on disk subsystem and its algorithms
- OS is responsible for the following activities:
 - Free-space management
 - Storage allocation
 - Disk scheduling
- Some storage need not be fast
 - Tertiary storage includes optical storage (CD, DVD), magnetic tape
 - Still must be managed – by OS or applications
 - Varies between WORM (write-once, read-many-times) and RW (read-write) formats



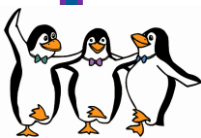


Storage Management

- Caching:

- Information is normally kept in main memory, when it is used, it is copied into a faster storage system (the cache) on a temporary basis.
- Cache management is an important design problem.
 - Cache size and the replacement policy can result in greatly increased performance
- Main memory can be viewed as a fast cache for secondary storage
- Performance of Various Levels of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape






Storage Management

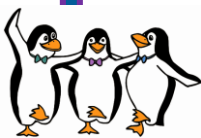
- Cache Coherency:

- In a hierarchical storage structure, the same data may appear in different levels of the storage system.
- For example

- Migration of integer “A” from Disk to Register



- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy
- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex
 - Several copies of a datum can be kept in different computers
 - Distributed systems must ensure that, when a replica is updated in one place, all other replicas are brought up to date on time

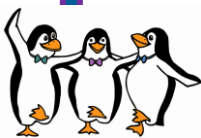




Storage Management

- I/O Systems:


- One purpose of OS is to hide peculiarities of hardware devices from the user
- The I/O subsystem consists of several components:
 - Memory management of I/O including **buffering** (storing data temporarily while it is being transferred), **caching** (storing parts of data in faster storage for performance), **spooling** (the overlapping of output of one job with input of other jobs)
 - General device-driver interface
 - Drivers for specific hardware devices
- Only the device driver knows the peculiarities of the specific device to which it is assigned.





Protection and Security


- Protection and Security:
 - **Protection:** any mechanism for controlling access of processes or users to resources defined by the OS
 - **Security:** defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
 - Systems generally first distinguish among users, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, one per user
 - User ID then associated with all files, processes of that user to determine access control
 - Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
 - **Privilege escalation** allows user to change to effective ID with more rights





Kernel Data Structures

- Lists, Stacks, and Queues:
 - An **array** is a simple data structure in which each element can be accessed directly.
 - For example, main memory is constructed as an array.
 - A **list** is the most fundamental data structures in computer science.
 - the items in a list must be accessed in a particular order.
 - **linked list** is the most common method for implementing the list,
 - Lists are used for constructing more powerful data structures, such as stacks and queues






Kernel Data Structures


- Lists, Stacks, and Queues:

- **Linked lists are of several types:**



- **Singly linked list**, each item points to its successor,



- **Doubly linked list**, a given item can refer either to its predecessor or to its successor



- **Circularly linked list**, the last element in the list refers to the first element, rather than to null,





Kernel Data Structures



- Lists, Stacks, and Queues:

- Stack:**

- It is a sequentially ordered data structure that uses the last in, first out (LIFO) principle for adding and removing items

- Queue:**

- It is a sequentially ordered data structure that uses the first in, first out (FIFO) principle








Kernel Data Structures

- Trees:

- A **tree** is a data structure that can be used to represent data hierarchically.
- Data values in a tree structure are linked through parent–child relationships
- In a **binary tree**, a parent may have at most two children (left child and the right child).
- A **binary search tree** additionally requires an ordering between the parent's two children in which $left_child \leq right_child$.
 - Search performance is $O(n)$



- An algorithm can be used to create a balanced binary search tree.
 - A tree containing **n** items has at most $\log_2(n)$ levels,
 - Search performance is $O(\log_2(n))$







Kernel Data Structures

- Hash Functions and Maps:

- A **hash function** takes data as its input, performs a numeric operation on this data, and returns a numeric value.
 - This numeric value can then be used as an index into a table (typically an array) to quickly retrieve the data.
 - Performance can be as good as $O(1)$
 - Hash functions are used extensively in operating systems.
 - Hash collision** can be accommodated by having a linked list at that table location that contains all of the items with the same hash value
 - A **hash map**, which associates (or maps) [key:value] pairs using a hash function

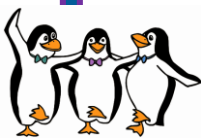




Kernel Data Structures

- Bitmaps:


- A **bitmap** is a string of **n** binary digits that can be used to represent the status of **n** items.
 - Example,
 - Consider the bitmap 001011101
 - Resources 2, 4, 5, 6, and 8 are unavailable; resources 0, 1, 3, and 7 are available.





Computing Environments


- Traditional Computing:
 - Stand-alone general purpose machines
 - But blurred as most systems interconnect with others (i.e., the Internet)
 - **Portals** provide web access to internal systems
 - **Network computers** (**thin clients**) are like Web terminals
 - Mobile computers interconnect via **wireless networks**
 - Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks





Computing Environments

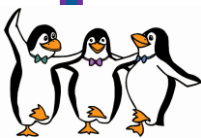
- Mobile Computing:
 - Handheld smartphones, tablets, etc
 - What is the functional difference between them and a “traditional” laptop?
 - Extra feature – more OS features (GPS, gyroscope)
 - Allows new types of apps like *augmented reality*
 - Use IEEE 802.11 wireless, or cellular data networks for connectivity
 - Leaders are **Apple iOS** and **Google Android**





Computing Environments

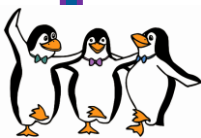

- Distributed Computing:
 - Distributed computing
 - Collection of separate, possibly heterogeneous, systems networked together
 - Network is a communications path, TCP/IP most common
 - Local Area Network (LAN)
 - Wide Area Network (WAN)
 - Metropolitan Area Network (MAN)
 - Personal Area Network (PAN)
 - Network Operating System provides features between systems across network
 - Communication scheme allows systems to exchange messages
 - Illusion of a single system





Computing Environments

- Client-Server:
 - Dumb terminals supplanted by smart PCs
 - Many systems now **servers**, responding to requests generated by **clients**
 - **Compute-server system** provides an interface to client to request services (i.e., database)
 - Example: A server running a database that responds to client requests for data
 - **File-server system** provides interface for clients to store and retrieve files
 - Example: a web server that delivers files to clients running web browsers.

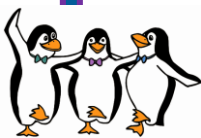





Computing Environments

- Peer-to-Peer:

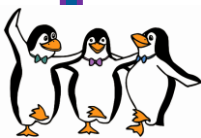
- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - Registers its service with central lookup service on network, or
 - Broadcast request for service and respond to requests for service via *discovery protocol*
- Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype





Computing Environments


- **Virtualization:**
 - Allows operating systems to run as applications within other OSes
 - Vast and growing industry
 - **Emulation** used when source CPU type different from target type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code – **Interpretation**
 - **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - **VMM** (virtual machine Manager) provides virtualization services





Computing Environments


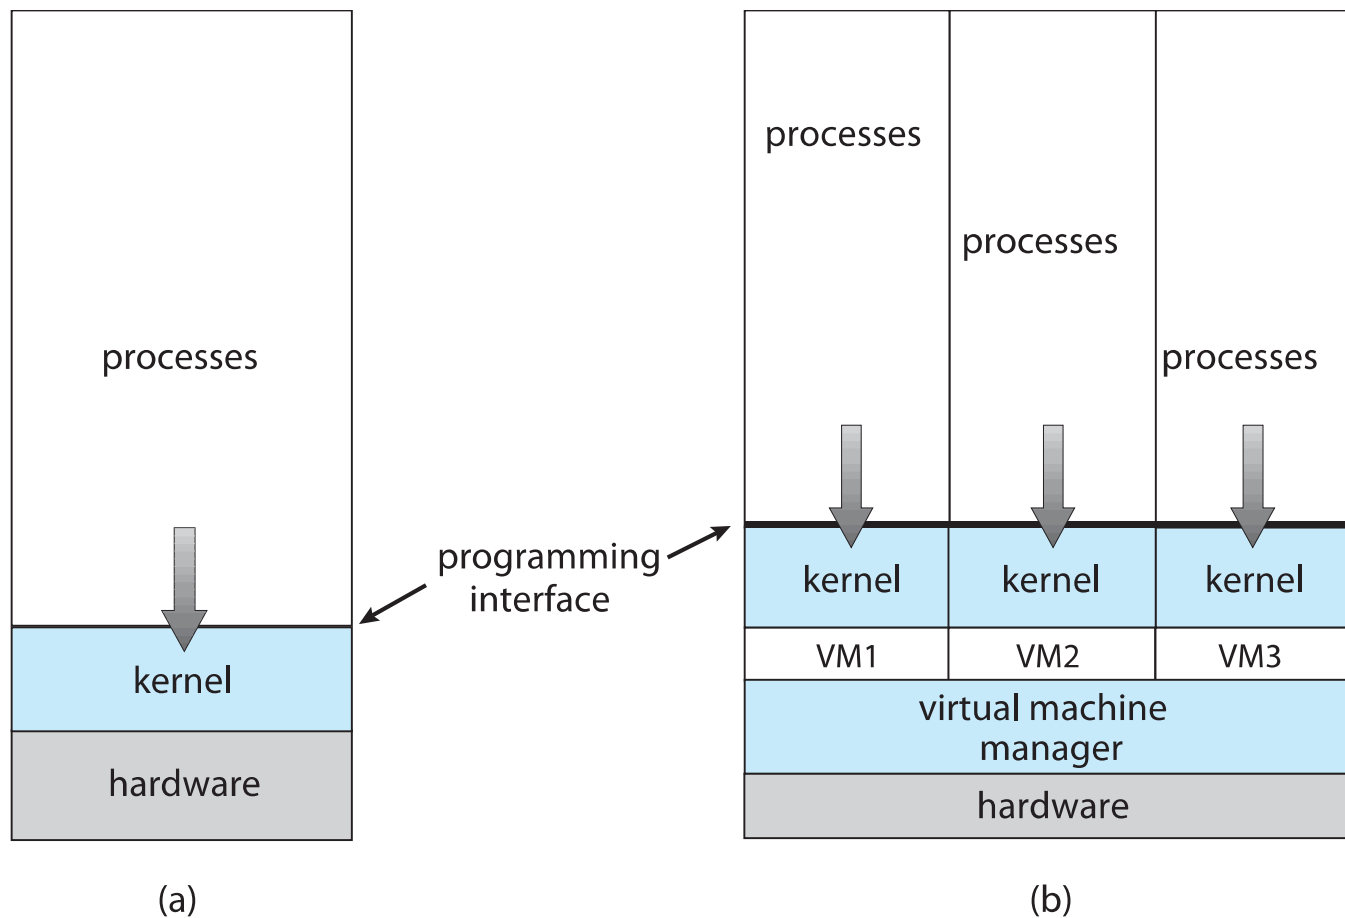
- Virtualization:
 - Use cases involve laptops and desktops running multiple OSES for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSES without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers
 - VMM can run natively, in which case they are also the host
 - There is no general purpose host then (VMware ESX and Citrix XenServer)





Computing Environments

- Virtualization:

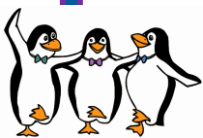




Computing Environments

- Cloud Computing:

- Delivers computing, storage, even apps as a service across a network
- It is logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Amazon Elastic Compute Cloud ([EC2](#)) has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage
- Many types of cloud computing:
 - [Public cloud](#) – available via Internet to anyone willing to pay
 - [Private cloud](#) – run by a company for the company's own use
 - [Hybrid cloud](#) – includes both public and private cloud components
 - Software as a Service ([SaaS](#)) – one or more applications available via the Internet (i.e., word processor)
 - Platform as a Service ([PaaS](#)) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service ([IaaS](#)) – servers or storage available over Internet (i.e., storage available for backup use)

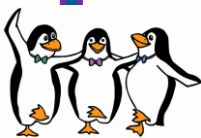





Computing Environments

- Cloud Computing:

- Cloud computing environments composed of traditional OSe, plus VMMs, plus cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications





Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary **closed-source**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration

